

# Aprendizaje Automático - Tarea 1

## Reconocimiento de letras manuscritas con SVM

Nicolas Farfán Cheneaux

```
import scipy.io
import pandas as pd
import random
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, accuracy_score,
classification_report
import numpy as np
from sklearn.svm import SVC
from scipy.io import loadmat
from collections import defaultdict
from sklearn.decomposition import PCA
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import reciprocal, uniform
```

### Lectura de datos

Dado que el dataset dado está en formato matlab (.mat) se lee mediante la biblioteca scipy

```
base_path = 'emnist-letters.mat'
data = loadmat(base_path)
```

Dada la documentación se logra obtener las muestras de imágenes y las etiquetas, indexando el arreglo data y almacenandolo en las variables imgs, labels. Para un mejor entendimiento se imprime los largos y valores de ejemplo para entender la estructura del dataset.

### Organización de los datos en diccionarios

Se coloca el total de imágenes por letra(clase) en un diccionario donde la key es cada label,

Se confirma que cada clase tiene 1000 muestras de entrenamiento y 100 de prueba (distintas entre si dado que se selecciona aleatoriamente 1100 muestras de cada clase)

### Organización de los datos de entrenamiento y testing

Dado que se tiene un total de 124.800 imágenes y un total de 26 letras, por cada clase se tendrá 4800 muestras. Se convierte los labels a numpy array y se cuentan las muestras por clase.

```
X_train_data = data['dataset']['train'][0, 0]['images'][0, 0]
y_train_data = data['dataset']['train'][0, 0]['labels'][0, 0]
```

```

X_test_data = data['dataset']['test'][0, 0]['images'][0, 0]
y_test_data = data['dataset']['test'][0, 0]['labels'][0, 0]

def organize_by_class(X, y):
    # Inicializa un diccionario para mantener los índices por clase
    class_dict = {i: [] for i in range(1, 27)} # Clases de 1 a 26
    for index, label in enumerate(y.ravel()): # Asegúrate de que 'y'
es un array 1D
        class_dict[label].append(index)
    return class_dict

# Organizar los datos de entrenamiento y prueba por clases
class_indices_train = organize_by_class(X_train_data, y_train_data)
class_indices_test = organize_by_class(X_test_data, y_test_data)

def select_samples(class_dict, num_samples):
    indices = []
    for label, idxs in class_dict.items():
        if len(idxs) < num_samples:
            raise ValueError(f"Clase {label} tiene menos ejemplos
({len(idxs)}) que el número de muestras deseado ({num_samples}).")
        selected_indices = np.random.choice(idxs, num_samples,
replace=False)
        indices.extend(selected_indices)
    return indices

# Seleccionar 1000 muestras de entrenamiento y 100 de prueba para cada
clase
train_indices = select_samples(class_indices_train, 1000)
test_indices = select_samples(class_indices_test, 100)

# Subconjunto de datos y etiquetas usando los índices aleatorios
X_train = X_train_data[train_indices]
y_train = y_train_data[train_indices].ravel() # Asegurarse de que y
es un arreglo 1D
X_test = X_test_data[test_indices]
y_test = y_test_data[test_indices].ravel() # Asegurarse de que y es
un arreglo 1D

```

## Experimento n°1

Se normalizan los vectores

```

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

## Entrenamiento del modelo SVM con kernel rbf

```
svm_model = SVC(kernel = "rbf")
svm_model.fit(X_train, y_train)

SVC()

param_distributions = {"gamma": reciprocal(0.001, 0.1), "C":
uniform(1, 10)}
rnd_search_cv = RandomizedSearchCV(svm_model, param_distributions,
n_iter=10, verbose=2, cv=3)
rnd_search_cv.fit(X_train[:3500], y_train.ravel()[:3500])

Fitting 3 folds for each of 10 candidates, totalling 30 fits
[CV] END ...C=6.300518934339031, gamma=0.0038021979868677292; total
time= 2.4s
[CV] END ...C=6.300518934339031, gamma=0.0038021979868677292; total
time= 2.3s
[CV] END ...C=6.300518934339031, gamma=0.0038021979868677292; total
time= 2.3s
[CV] END ....C=5.382781342365406, gamma=0.011198469181398213; total
time= 3.7s
[CV] END ....C=5.382781342365406, gamma=0.011198469181398213; total
time= 4.1s
[CV] END ....C=5.382781342365406, gamma=0.011198469181398213; total
time= 3.7s
[CV] END ....C=2.72273382257199, gamma=0.0011074830352197256; total
time= 1.2s
[CV] END ....C=2.72273382257199, gamma=0.0011074830352197256; total
time= 1.4s
[CV] END ....C=2.72273382257199, gamma=0.0011074830352197256; total
time= 1.3s
[CV] END ...C=4.989449697222545, gamma=0.0011107693145562913; total
time= 1.2s
[CV] END ...C=4.989449697222545, gamma=0.0011107693145562913; total
time= 1.2s
[CV] END ...C=4.989449697222545, gamma=0.0011107693145562913; total
time= 1.2s
[CV] END ..C=10.017909601866481, gamma=0.0039771858299884385; total
time= 2.4s
[CV] END ..C=10.017909601866481, gamma=0.0039771858299884385; total
time= 2.5s
[CV] END ..C=10.017909601866481, gamma=0.0039771858299884385; total
time= 2.5s
[CV] END ....C=9.357365930430243, gamma=0.001029618789020042; total
time= 1.2s
[CV] END ....C=9.357365930430243, gamma=0.001029618789020042; total
time= 1.2s
[CV] END ....C=9.357365930430243, gamma=0.001029618789020042; total
time= 1.2s
[CV] END ...C=3.7635650684612134, gamma=0.011063129174722887; total
```

```

time= 3.6s
[CV] END ...C=3.7635650684612134, gamma=0.011063129174722887; total
time= 3.6s
[CV] END ...C=3.7635650684612134, gamma=0.011063129174722887; total
time= 3.6s
[CV] END .....C=7.612975768275725, gamma=0.0262580189994162; total
time= 3.7s
[CV] END .....C=7.612975768275725, gamma=0.0262580189994162; total
time= 3.9s
[CV] END .....C=7.612975768275725, gamma=0.0262580189994162; total
time= 3.8s
[CV] END ....C=2.959942989849772, gamma=0.004415724708738391; total
time= 2.8s
[CV] END ....C=2.959942989849772, gamma=0.004415724708738391; total
time= 2.7s
[CV] END ....C=2.959942989849772, gamma=0.004415724708738391; total
time= 2.7s
[CV] END ..C=5.3694547336146305, gamma=0.0010648018753419805; total
time= 1.2s
[CV] END ..C=5.3694547336146305, gamma=0.0010648018753419805; total
time= 1.1s
[CV] END ..C=5.3694547336146305, gamma=0.0010648018753419805; total
time= 1.2s

```

```

RandomizedSearchCV(cv=3, estimator=SVC(),
                   param_distributions={'C':
<scipy.stats._distr_infrastructure.rv_continuous_frozen object at
0x000002B411945C10>,
                                     'gamma':
<scipy.stats._distr_infrastructure.rv_continuous_frozen object at
0x000002B4118F83E0>},
                   verbose=2)

```

```

rnd_search_cv.best_estimator_

```

```

SVC(C=9.357365930430243, gamma=0.001029618789020042)

```

```

rnd_search_cv.best_score_

```

```

0.945713378632814

```

```

svm_model_adjusted = SVC(kernel = "rbf", C = 9.357365930430243, gamma
= 0.001029618789020042)
svm_model_adjusted.fit(X_train, y_train)

```

```

SVC(C=9.357365930430243, gamma=0.001029618789020042)

```

## Accuracy

Función **get\_accuracy**, se usa la función predict para obtener las etiquetas predichas por el modelo y usando la matriz de confusión se calcula la precisión de cada categoría

```
def get_accuracy(X_test, y_test, model):
    # Predicciones del modelo, dado las imágenes de prueba se obtiene
    los labels
    y_pred = model.predict(X_test)

    # Matriz de confusión (Resultados reales vs predicciones)
    cm = confusion_matrix(y_test, y_pred)

    # Calcular la precisión por categoría
    accuracies = cm.diagonal() / cm.sum(axis=1)
    print(classification_report(y_test, y_pred, zero_division=0))
```

Accuracy por Categoría y Total de modelo con y sin ajuste de hiperparámetros

Se muestra las accuracies por categoría de la letra "A"(1) a la "Z"(26). Y además la accuracy total 87%

```
get_accuracy(X_test, y_test, svm_model)
```

	precision	recall	f1-score	support
1	0.77	0.85	0.81	100
2	0.89	0.91	0.90	100
3	0.89	0.90	0.90	100
4	0.88	0.81	0.84	100
5	0.93	0.85	0.89	100
6	0.88	0.79	0.83	100
7	0.72	0.68	0.70	100
8	0.87	0.78	0.82	100
9	0.67	0.72	0.69	100
10	0.89	0.93	0.91	100
11	0.90	0.89	0.89	100
12	0.69	0.66	0.68	100
13	0.91	0.90	0.90	100
14	0.85	0.85	0.85	100
15	0.87	0.96	0.91	100
16	0.86	0.95	0.90	100
17	0.73	0.76	0.75	100
18	0.84	0.87	0.85	100
19	0.99	0.93	0.96	100
20	0.87	0.88	0.88	100
21	0.94	0.94	0.94	100
22	0.95	0.86	0.90	100
23	0.69	0.94	0.79	100
24	0.92	0.88	0.90	100
25	0.92	0.81	0.86	100
26	0.98	0.86	0.91	100
accuracy			0.85	2600
macro avg	0.86	0.85	0.85	2600

weighted avg	0.86	0.85	0.85	2600
--------------	------	------	------	------

```
get_accuracy(X_test, y_test, svm_model_adjusted)
```

	precision	recall	f1-score	support
1	0.77	0.86	0.82	100
2	0.93	0.95	0.94	100
3	0.87	0.92	0.89	100
4	0.94	0.85	0.89	100
5	0.91	0.87	0.89	100
6	0.85	0.82	0.84	100
7	0.75	0.74	0.74	100
8	0.93	0.83	0.88	100
9	0.69	0.71	0.70	100
10	0.90	0.92	0.91	100
11	0.89	0.91	0.90	100
12	0.66	0.67	0.66	100
13	0.89	0.92	0.91	100
14	0.92	0.84	0.88	100
15	0.91	0.98	0.94	100
16	0.88	0.96	0.92	100
17	0.77	0.72	0.74	100
18	0.84	0.87	0.86	100
19	0.99	0.95	0.97	100
20	0.89	0.85	0.87	100
21	0.96	0.90	0.93	100
22	0.90	0.90	0.90	100
23	0.73	0.94	0.82	100
24	0.93	0.90	0.91	100
25	0.89	0.85	0.87	100
26	0.98	0.87	0.92	100
accuracy			0.87	2600
macro avg	0.87	0.87	0.87	2600
weighted avg	0.87	0.87	0.87	2600

Visualización de un ejemplo con la letra real vs la predicción

```
def visualize_prediction(X, y, model, index):  
    image = X[index].reshape(28, 28) # Redimensionar el vector a  
    28x28 para visualización  
    image = np.rot90(image, k=-1)  
    image = np.fliplr(image)  
    true_label_index = y[index] # Índice de la etiqueta verdadera  
    predicted_label_index = model.predict([X[index]])[0]  
  
    true_label = chr(true_label_index + 64) # 64 + 1 = 65, que es 'A'
```

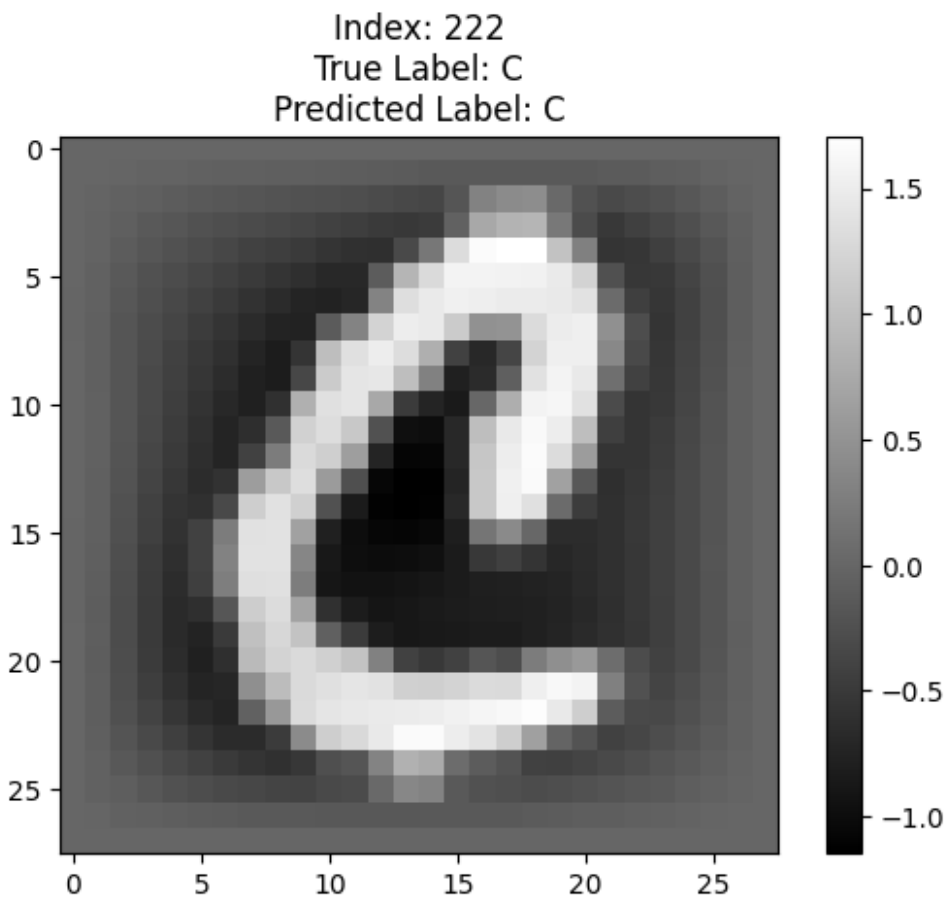
```

en Unicode
predicted_label = chr(predicted_label_index + 64) # 64 + 'índice'
para obtener la letra correcta

# Mostrar la imagen
plt.imshow(image, cmap='gray')
plt.title(f'Index: {index}\nTrue Label: {true_label}\nPredicted
Label: {predicted_label}')
plt.colorbar()
plt.grid(False)
plt.show()

visualize_prediction(X_test, y_test, svm_model, 222)

```



## Experimento n°2

Se utiliza PCA para reducir de 784 a 128 la dimensión de los vectores.

```

# Inicializar PCA con 128 componentes
pca = PCA(n_components = 128)

```

```
# Entrena PCA
pca.fit(X_train)

# Ajustar y transformar los datos
X_train_pca = pca.transform(X_train)
X_test_pca = pca.transform(X_test)
X_train_pca.shape

(26000, 128)
```

Entrenamiento del modelo con los datos reducidos por PCA

```
svm_model_pca = SVC(kernel='rbf')
svm_model_pca.fit(X_train_pca, y_train)

SVC()
```

Accuracy por Categoría y Total

```
param_distributions = {"gamma": reciprocal(0.001, 0.1), "C":
uniform(1, 10)}
rnd_search_pca = RandomizedSearchCV(svm_model_pca,
param_distributions, n_iter=10, verbose=2, cv=3)
rnd_search_pca.fit(X_train_pca[:3500], y_train.ravel()[:3500])

Fitting 3 folds for each of 10 candidates, totalling 30 fits
[CV] END ....C=9.593409474005007, gamma=0.004646228193086938; total
time= 0.5s
[CV] END ....C=9.593409474005007, gamma=0.004646228193086938; total
time= 0.5s
[CV] END ....C=9.593409474005007, gamma=0.004646228193086938; total
time= 0.5s
[CV] END .....C=7.06314647947614, gamma=0.004157162830130484; total
time= 0.5s
[CV] END .....C=7.06314647947614, gamma=0.004157162830130484; total
time= 0.5s
[CV] END .....C=7.06314647947614, gamma=0.004157162830130484; total
time= 0.5s
[CV] END ...C=2.4254956523831352, gamma=0.056862375824384914; total
time= 0.9s
[CV] END ...C=2.4254956523831352, gamma=0.056862375824384914; total
time= 0.9s
[CV] END ...C=2.4254956523831352, gamma=0.056862375824384914; total
time= 0.9s
[CV] END ....C=6.801719455631144, gamma=0.012604102266671541; total
time= 0.8s
[CV] END ....C=6.801719455631144, gamma=0.012604102266671541; total
time= 0.8s
[CV] END ....C=6.801719455631144, gamma=0.012604102266671541; total
time= 0.8s
```



```

[CV] END .....C=8.383218276759134, gamma=0.0468218134762736; total
time= 0.9s
[CV] END .....C=8.383218276759134, gamma=0.0468218134762736; total
time= 0.9s
[CV] END .....C=8.383218276759134, gamma=0.0468218134762736; total
time= 0.9s
[CV] END ....C=1.4736728486408692, gamma=0.06200894731251944; total
time= 0.9s
[CV] END ....C=1.4736728486408692, gamma=0.06200894731251944; total
time= 0.9s
[CV] END ....C=1.4736728486408692, gamma=0.06200894731251944; total
time= 0.9s
[CV] END .....C=9.740406647771055, gamma=0.06180389065378497; total
time= 0.9s
[CV] END .....C=9.740406647771055, gamma=0.06180389065378497; total
time= 0.9s
[CV] END .....C=9.740406647771055, gamma=0.06180389065378497; total
time= 0.9s
[CV] END ...C=6.966685478321949, gamma=0.0014151628625696535; total
time= 0.2s
[CV] END ...C=6.966685478321949, gamma=0.0014151628625696535; total
time= 0.2s
[CV] END ...C=6.966685478321949, gamma=0.0014151628625696535; total
time= 0.2s
[CV] END ....C=5.296262803728709, gamma=0.002336143805003388; total
time= 0.3s
[CV] END ....C=5.296262803728709, gamma=0.002336143805003388; total
time= 0.3s
[CV] END ....C=5.296262803728709, gamma=0.002336143805003388; total
time= 0.3s
[CV] END .....C=9.485177062636275, gamma=0.05576765804432435; total
time= 0.9s
[CV] END .....C=9.485177062636275, gamma=0.05576765804432435; total
time= 0.9s
[CV] END .....C=9.485177062636275, gamma=0.05576765804432435; total
time= 0.9s

```

```

RandomizedSearchCV(cv=3, estimator=SVC(),
                    param_distributions={'C':
<scipy.stats._distn_infrastructure.rv_continuous_frozen object at
0x000002B4118D7A10>,
                                        'gamma':
<scipy.stats._distn_infrastructure.rv_continuous_frozen object at
0x000002B414336B40>},
                    verbose=2)

```

```

best_svmpca_model = rnd_search_pca.best_estimator_
best_svmpca_model

```

```

SVC(C=6.966685478321949, gamma=0.0014151628625696535)

```

```
rnd_search_cv.best_score_
```

```
0.945713378632814
```

```
svm_modelpca_adjusted = SVC(kernel = "rbf", C = 6.966685478321949,  
gamma=0.0014151628625696535)
```

```
svm_modelpca_adjusted.fit(X_train_pca, y_train)
```

```
SVC(C=6.966685478321949, gamma=0.0014151628625696535)
```

```
get_accuracy(X_test_pca, y_test, svm_model_pca)
```

	precision	recall	f1-score	support
1	0.79	0.85	0.82	100
2	0.93	0.94	0.94	100
3	0.90	0.91	0.91	100
4	0.90	0.83	0.86	100
5	0.95	0.87	0.91	100
6	0.89	0.81	0.85	100
7	0.76	0.73	0.74	100
8	0.87	0.81	0.84	100
9	0.68	0.70	0.69	100
10	0.91	0.94	0.93	100
11	0.89	0.90	0.90	100
12	0.68	0.68	0.68	100
13	0.91	0.92	0.92	100
14	0.89	0.86	0.87	100
15	0.88	0.98	0.92	100
16	0.85	0.94	0.90	100
17	0.75	0.77	0.76	100
18	0.86	0.87	0.87	100
19	0.99	0.94	0.96	100
20	0.88	0.90	0.89	100
21	0.93	0.94	0.94	100
22	0.95	0.86	0.90	100
23	0.76	0.94	0.84	100
24	0.91	0.91	0.91	100
25	0.91	0.83	0.87	100
26	0.98	0.90	0.94	100
accuracy			0.87	2600
macro avg	0.87	0.87	0.87	2600
weighted avg	0.87	0.87	0.87	2600

```
get_accuracy(X_test_pca, y_test, svm_modelpca_adjusted)
```

	precision	recall	f1-score	support
1	0.81	0.88	0.85	100

2	0.90	0.95	0.93	100
3	0.90	0.92	0.91	100
4	0.95	0.89	0.92	100
5	0.91	0.91	0.91	100
6	0.93	0.85	0.89	100
7	0.75	0.72	0.73	100
8	0.90	0.85	0.88	100
9	0.68	0.73	0.71	100
10	0.93	0.94	0.94	100
11	0.88	0.92	0.90	100
12	0.69	0.66	0.67	100
13	0.92	0.92	0.92	100
14	0.91	0.88	0.89	100
15	0.90	0.98	0.94	100
16	0.92	0.97	0.94	100
17	0.78	0.73	0.76	100
18	0.86	0.89	0.87	100
19	0.98	0.96	0.97	100
20	0.92	0.89	0.90	100
21	0.96	0.91	0.93	100
22	0.91	0.88	0.89	100
23	0.79	0.95	0.86	100
24	0.93	0.91	0.92	100
25	0.92	0.86	0.89	100
26	0.97	0.91	0.94	100
accuracy			0.88	2600
macro avg	0.88	0.88	0.88	2600
weighted avg	0.88	0.88	0.88	2600

## Experimento nº3

Vectores de características entrenados extraídos de una red convolucional simple entrenada en EMNIST

```
# Dataframe de entrenamiento
df_train = pd.read_csv('train_embeddings.csv')
df_train = df_train.sample(n=1000, random_state=42)

# Dataframe de prueba
df_test = pd.read_csv('test_embeddings.csv')
df_test = df_test.sample(n=100, random_state=42)

X_train = df_train.drop('label', axis=1).values
y_train = df_train['label'].values

# Se divide la data en características y labels
```

```
X_test = df_test.drop('label', axis=1).values
y_test = df_test['label'].values
```

Entrenamiento del modelo con los vectores de característica dados

```
svm_vec = SVC(kernel = "rbf")
svm_vec.fit(X_train, y_train)

SVC()
```

Accuracy por Categoría y Total

```
get_accuracy(X_test, y_test, svm_vec)
```

	precision	recall	f1-score	support
0	0.67	1.00	0.80	2
1	1.00	0.67	0.80	3
2	1.00	1.00	1.00	4
3	0.00	0.00	0.00	0
4	1.00	0.80	0.89	5
5	1.00	0.50	0.67	2
6	0.67	0.67	0.67	3
7	1.00	1.00	1.00	3
8	0.57	1.00	0.73	4
9	1.00	0.50	0.67	2
10	1.00	1.00	1.00	5
11	1.00	0.60	0.75	5
12	1.00	1.00	1.00	5
13	1.00	1.00	1.00	4
14	1.00	1.00	1.00	3
15	0.86	1.00	0.92	6
16	0.86	0.86	0.86	7
17	1.00	1.00	1.00	5
18	0.83	0.83	0.83	6
19	0.67	1.00	0.80	2
20	1.00	1.00	1.00	5
21	1.00	1.00	1.00	4
22	1.00	1.00	1.00	2
23	1.00	1.00	1.00	6
24	1.00	0.67	0.80	3
25	1.00	1.00	1.00	4
accuracy				100
macro avg				0.89 0.85 0.85 100
weighted avg				0.93 0.90 0.90 100

```
C:\Users\nlc8l\AppData\Local\Temp\ipykernel_18064\1174503523.py:9:  
RuntimeWarning: invalid value encountered in divide  
    accuracies = cm.diagonal() / cm.sum(axis=1)
```